

- Continuous Integration (CI)
  - ↳ Developers continuously integrate (merge) their code changes into a shared branch (like main) & every time they do, the system automatically builds, test, & verifies the code
  - ↳ GitHub Actions / Jenkins / GitLab CI pipelines automatically runs
    - i - Unit Tests
    - ii - Code Linting
    - iii - Build Checks
- ↳ If all pass → code is integrated successfully

### → Continuous Delivery (CD - Delivery)

Once code passes all tests (CI), it is automatically packaged & deployed to a staging or testing env ready for release to production — but the final step (production deployment) is manual

### → Continuous Deployment (CD - Deployment)

One step further — every change that passes thru (CI/CD) pipeline is automatically deployed to production — fully automated (no approval)

Date \_\_\_\_\_

Day \_\_\_\_\_

"Remove the handoffs - streamline the process, challenge everything that doesn't add value"

→ CAMS (Culture Automation Measurement Sharing)

Culture — Collaboration, trust and communication between Dev, Ops, QA & other teams

- ↳ Break down silos and promote shared responsibility for SLAs delivery
- "People over processes"

Automation — Automate repetitive & error-prone tasks to increase speed & reliability

- ↳ CI/CD pipelines, Automated Testing, monitoring & deployments

Lean — Focus on efficiency, eliminate waste and continuously improve

- ↳ Smaller frequent releases, instead of big, risky ones

↳ Short feedback loops b/w dev & users

Measurement — Encourage open communications, shared learnings & transparency

- ↳ Documenting best practices

Track key metrics to understand performance & find areas of improvement

Date \_\_\_\_\_

Day \_\_\_\_\_

Starting

- Encourage open communication,  
Shared learnings & transparency
- ↳ Documenting best practices
- ↳ Sharing incident reports

→ Should we have a DevOps team or hire DevOps Engineers?

↳ DevOps is not a job title — It's a culture and set of practices.

↳ The goal of DevOps is to break the wall between Development and Operations, not to create a new wall called DevOps team.

↳ If you are following CALMS, then you are already doing DevOps

↳ If you create another team that become yet another silo

→ DevOps Principles : Three ways

① System Thinking

↳ Optimize the entire system, not just one team or department

↳ The goal is to make workflow efficient from code creation to deployment

↳ Identify & remove bottlenecks.

Date \_\_\_\_\_

Day \_\_\_\_\_

## ② Amplifying Feedback Loop

- ↳ Create short & continuous feedback loops between developers, testers, operations & customers.
- ↳ Detect problem early, learn quickly
- ↳ Continuous monitoring & alerting
- ↳ Build automated tests into the pipeline so developers can get feedback early

## ③ A culture of Continual Experimentation & learning

- ↳ Encourage risk taking & failing forward
- ↳ Allocating time for improvement work
- ↳ Rewarding teams for taking risk
- ↳ Encourage to learn from failures
- ↳ Invest in training

## ▷ Seven Principles & Seven Wastes of Lean

### ① Eliminating waste

- ↳ Not code more features than needed
- ↳ Minimize handoffs
- ↳ Take right decisions at right time

### ② Building quality

### ③ Creating knowledge

### ④ Differing commitment

Date \_\_\_\_\_

Day \_\_\_\_\_

- ⑤ Deliver fast
- ⑥ Respect People
- ⑦ Optimize the whole

↳ System thinking approach

→ use to create culture of continuous improvement & continuous learning

▷ A3 Problem Solving Framework

- ① Set the background

↳ include stat of how problem directly impact business outcomes

- ② Current Condition & Problem Statement

- ③ Develop Goal

- ④ Performance Analysis

- ⑤ Brainstorm

- ⑥ Countermeasure

- ⑦ Implementation plan

- ⑧ Update → update "std work" based on above steps

Date \_\_\_\_\_

Day \_\_\_\_\_

## "INTRODUCTION TO AWS"

### ▷ On-Premise

- ↳ Everything managed by your own company - servers, storage, network, software
- ↳ A bank running its own datacenter

### ▷ IaaS

- ↳ Infrastructure as a Service
- ↳ Cloud provider basic infra - servers, storage and networking - you manage OS & apps
- ↳ Amazon EC2, Google Compute Engine
- ↳ You rent VMs & install OS or any software that is needed

### ▷ PaaS

- ↳ Platform as a Service
- ↳ Cloud provide ready made platform to build, run, deploy
- ↳ Heroku, elastic beanstalk

### ▷ SaaS

- ↳ Software as a Service
- ↳ Fully functional software
- ↳ Gmail, Google Docs, Dropbox

Date \_\_\_\_\_

Day \_\_\_\_\_

- Regions :

- A region is a geographical area in the world where AWS has multiple data centers

- ↳ Choose region close to the intended users for low latency & to comply with local laws.

- Availability Zones (AZs)

- Each region has multiple availability zones (usually 3 or more)

- Each AZ is basically a separate data center (cluster of data center) - with independent power, networking & cooling

- Connected to each other with high speed & low-latency links

High Availability → Within a region across diff AZs

Disaster Recovery → Across Region

- Local Zones

- Brings AWS compute, storage and database services closer to end users in specific cities (outside region)

- They extend the region to a specific metro area for ultra-low-latency

Page No. [cgs gaming, streaming]

## o Wavelength Zones

- These are AWS infra deployed inside telecom networks (5G Networks).
- They bring compute & storage very close to mobile devices.
- Mobile Gaming, AR/VR, autonomous vehicle etc.

## o AWS Outposts

- Outposts brings AWS infra physically into your own on-premises data center.
- A hospital keeping sensitive patient data locally but using Outposts for compute.

## ▷ Selecting AWS Region

### (1) Proximity to Users

↳ Select region where most of users are located

### (2) Compliance

↳ Some countries have laws that require data to stay within nation boundaries

↳ Choose region which meet regulations

### (3) Service Availability

↳ Not all AWS services available in every region

### (4) Cost

↳ Prices varies by region due to energy, land and tax differences

Date \_\_\_\_\_

Day \_\_\_\_\_

## " AMAZON-VPC "

- Amazon Virtual Private Cloud (VPC) allows user to use AWS resources in a virtual network.
- The users can customize virtual networking environments like
  - ↳ Selecting own IP range
  - ↳ Creating subnets
  - ↳ Configuring route tables & network gateways

### ► Core Components

- VPC CIDR Block
- Subnet
- Gateways
- Route Table
- Network Access Control Lists (NACLs)
- Security Groups

- VPC is a regional service

Date \_\_\_\_\_

Day \_\_\_\_\_

## ► Features

- Specify own private IP ranges of your choice
- Expand VPC by adding secondary IP ranges
  - 10.0.0.0/16 → Primary
  - 192.168.0.0/24 → Secondary
- Divide VPC private IP Address range into one or more public or private subnets
- Control inbound & outbound access using NACLs
- Store Data in S3 with set permissions such that data can only be accessed from within VPC
- Attach a public IP to any instance in VPC so it is accessible directly from Internet
- Connect multiple VPCs thru VPC peering or transit gateway
- Privately connect other AWS services without using Internet, thru NAT or VPC endpoints
- Bridge VPC & onsite IT infra with an encrypted VPN connection
- VPC Flow logs for monitoring

Date \_\_\_\_\_

Day \_\_\_\_\_

## ▷ Subnets

- Range of IP Addrs in VPC
- Using public subnet for resources that must be connected to Internet gateway
- And private subnet that should not be accessible
- Each subnet must reside with only one AZs
- Each AZs can have multiple diff subnets
- An Internet Gateway (igw) enables communication over internet
  - ↳ A subnet attached to igw is called public subnet
  - ↳ An instance must have a public IPv4 or an Elastic IP Address to communicate with internet
  - ↳ A VPC can only be connected with one & only one igw.
- A subnet that has no route to igw is called a private subnet
- A Virtual Private Network (VPN) connection enables communication with your corporate network.

- Subnet Sizing
  - ↳ The allowed block size is below a /16 netmask (65,536 IP Addresses) and a /28 netmask (16 IP Addresses).
  - ↳ You can also associate a Secondary VPC.
- Total 5 IP's are reserved for each subnet
  - ↳ The 1<sup>st</sup>, four & last.
  - ↳ In a subnet with CIDR block 10.0.0.0/24 the following IPs will be reserved
    - i - 10.0.0.0 : Network Address
    - ii - 10.0.0.1 : VPC Router (Gateway)
    - iii - 10.0.0.2 : Amazon DNS Server to allow for name resolutions throughout the network
    - iv - 10.0.0.3 : Reserved by AWS for future use
    - v - 10.0.0.255 : Network Broadcast Address
- VPC Endpoints
  - ↳ It lets thru thru diff AWS services to communicate without using IGW, NAT communicate b/w VPC & outside world.
  - ↳ Instances in a VPC do not req public IP to communicate with resources in the same service
  - ↳ Traffic remains in Amazon network

Date \_\_\_\_\_

Day \_\_\_\_\_

### ↳ Benefits:

- ↳ No internet exposure (safe)
- ↳ no need for NAT (cost saving)
- ↳ stays on AWS backbone network (low latency)

### Example:

- ↳ Let say you have:
- ↳ EC2 instances in a private network
- ↳ Your app uploads data to S3

Without VPC

NAT Gateway → Internet → AWS S3

With VPC

EC2 → Private Route → AWS S3

(Traffic stays inside AWS)

### ▷ NAT Gateway

- ↳ A managed AWS service that allows private subnet instances (without public IPs) to access the Internet but blocks incoming connections from the Internet
- ↳ Outbound Internet access only
- ↳ Used by instances to download updates or call external API's

PVT EC2 → NAT Gateway → IGW

Date \_\_\_\_\_

Day \_\_\_\_\_

Security Groups	NACLs
① Operates at instance level	① Operates at subnet level
② All inbound traffic is denied, all outbound allowed (Default)	② All outbound & inbound allowed (Default)
③ Only Allow Rule	③ Both Allow & Deny rules
④ All rules are evaluated together; the most permissive one applies	④ Rules are evaluated in order (lowest → highest) until a match is found
⑤ Applies to individual instances	⑤ Applies to all resources within a subnet
⑥ Act as firewall at instance level	⑥ Act as a firewall at subnet level
⑦ Stateful: Return traffic automatically allowed	⑦ Stateless: Return traffic must be explicitly allowed by rule

#### ▶ Flow Logs

↳ Capture info about IP traffic going to & from network interfaces in VPC

Date \_\_\_\_\_

Day \_\_\_\_\_

## ▷ VPC Peering

→ Allows two VPCs to communicate with each other using private IP addresses, as if they were part of same network.

→ It is like building a secure tunnel b/w two buildings (VPCs)

→ Bidirectional connection

→ Can be within region or across

→ Must update route tables in both VPCs to send traffic thru peering

→ No transitive property

VPC A → VPC B

VPC B → VPC C

VPC A → VPC C X

→ Cannot have overlapping CIDR blocks b/w VPCs

→ Here the one initiating the connection is the requestor & other is acceptor.

VPC A (10.0.0.0/16)

VPC B (192.168.0.0/16)

[EC2-1] ← Peering → [EC2-2]

(Private IP 10.0.1.8)

(Private IP  
192.168.1.2)

Date \_\_\_\_\_

Day \_\_\_\_\_

→ Edge to edge routing through a gateway or put network not allowed. As it would become transitive

(A) → (peering) → B → (Internet Gateway)

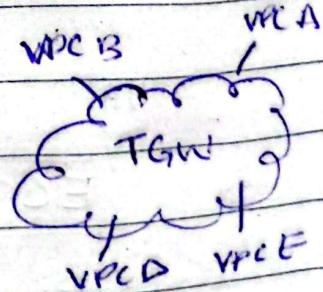
↳ ~~Since all VPCs are isolated so it~~

AWS intentionally forbids this to preserve network isolation & security b/w VPCs

### ▷ Transit Gateway (TGW)

→ A central hub that connects multiple VPCs, on-prem networks and even other AWS regions — all thru a single, scalable gateway  
→ Connect everything once to the TGW — it routes traffic intelligently

→ In Peering you would need  $\frac{n(n-1)}{2}$  connections whereas in TGW only "n" connections needed  
→ Transitive routing supported  
→ Can control which VPC talk to which one



Date \_\_\_\_\_

Day \_\_\_\_\_

## " AMAZON - EC2 "

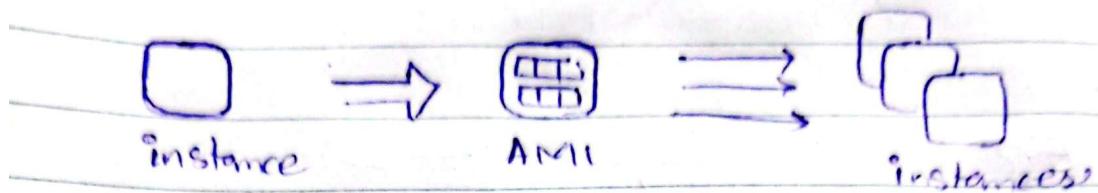
- Elastic Cloud Compute — allows to run applications & programs in the cloud env"
- Enables to vary the capacity within minutes
- EC2 virtual computing env, known as instances.
- Preconfigured templates — Amazon Machine Images (AMI)
- Persistent storage volumes for the data using Amazon Elastic Block Store (Amazon EBS volume)
- Also have security groups
- Temporary data storage deleted when instance stop or terminate known as instance store volumes

### ▷ Amazon Machine Image

" Provides the info required to launch an instance, which is a virtual server in the cloud."

→ You must specify src AMI when launch an instance

→ A block device mapping specifies the volumes to attach to the instance when it's launched



- Amazon Elastic Block Store (EBS)

→ Provides persistent, highly available, consistent low-latency block storage volumes for use with EC2

→ Block Storage service for EC2

→ Works as virtual hard drive

→ Each EBS volume exists in one AZ (not across region). Replicated automatically with the same AZ for fault tolerance.

→ Provides encryption

→ Protect data by creating point-in-time snapshots of EBS volumes.

→ Access Management

IOPS : No. of read & write operations per second

Throughput : No. of bits read or write per second

▷ GB v GiB

GB	GigaByte	Decimal (Base 10)	$1\text{GB} = 10^9$
GiB	GibiByte	Binary (Base 2)	$1\text{GiB} = 2^{30}$

Date \_\_\_\_\_

Day \_\_\_\_\_

$$100 \text{ GiB} = 100 \times 1.073741824 \text{ GB} \\ = 107.37 \text{ GB}$$

### ▷ Stop Vs Terminate

↓      ↓

Instance      instance permanently  
is shutdown      deleted along with  
but memory      EBS.  
remains intact

can be restarted      Cannot be restarted