

KUBERNETES

Date _____

- Open source, container orchestration tool
- Helps you manage containerized applications in different environments

The need for a container orchestration tool

- Trend from Monolithic to microservice architecture
- Increase usage of containers
- Demand for a proper way of managing those hundreds of containers

Features

- ↳ High availability / No downtime
 - ↳ Scalability or high performance
 - ↳ Disaster recovery - backup and restore
-
- Through Kubernetes we can run our app on any cloud provider. It provides an abstraction layer for us.
 - Immutable Infra → For making any change to server → Create copy of server from base image → changes made → replace old with new one

Date

Main Components of Kubernetes

• Pod

↳ Smallest unit of K8s

↳ Abstraction over containers

↳ It contains one or more application containers that are tightly coupled

↳ But mostly we should run one-container per Pod

↳ It creates a layer or running env on top of container

↳ K8s has made this abstraction so that we are not dependent on ~~contain~~ a specific containerization technology.

↳ Each Pod gets its own IP Addr (not the container!)

↳ Pods can communicate ~~with~~ with each other using IP

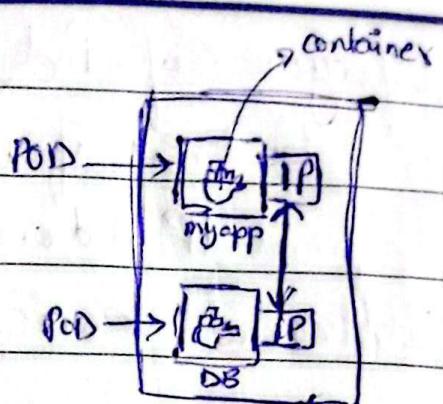
↳ Pods can die easily, so when they created again they get new IP. So if two pods communicate each other then we should update the IP.

• Service

↳ If we make a DB pod then if we make new DB, it will get new IP and we will have to change the IP in our application

↳ For this reason service is used

↳ It is basically a permanent IP Addr that can be attached to each Pod-



→ Service is also a load balancer for Pod. like we have two Pods of AddToCart Service. So the service sends request to the ideal Pod.

Date _____

↳ Lifecycle of Pod & service not connected.

↳ Even if Pod dies it will not loose its permanent IP Adr.

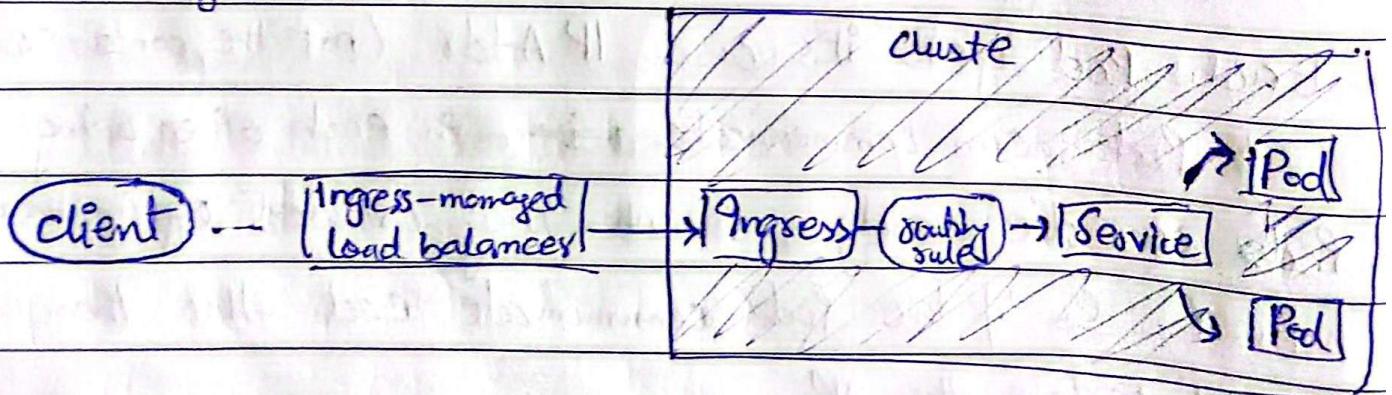
→ External Service (like we need to run an app on browser (easily accessible))

→ Internal Service (like DBs)

Ingress

→ Exposes HTTP & HTTPS routes from outside the cluster to services within the cluster.

→ Traffic routing is controlled by rules defined on the ingress resource.



- ConfigMap

- ↳ Used to store non-confidential data in key value pairs
- ↳ Pods can use them as env variables.
- ↳ ConfigMap allows you to decouple env-specific configuration from your container images
- ↳ ConfigMaps can also be mounted as data volumes

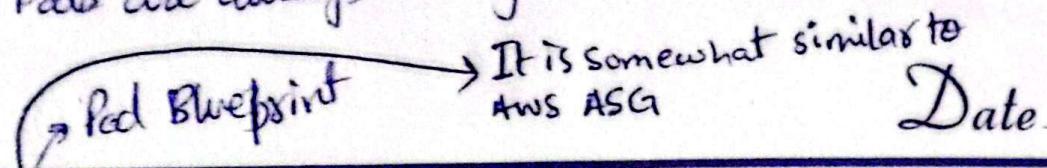
- Secret

- ↳ It contains small amount of sensitive data such as a password / a token or a key.
- ↳ They are similar to ConfigMap.
- ↳ They store data in base64 encoding format

- Volumes

- ↳ By default there is no data persistence on restarting Pod.
- ↳ But can be done ~~not~~ thru volume
- ↳ This volume is separate from cluster.
- ↳ It can be the part of the local machine or remote

ReplicaSet: Is a Kubernetes controller that ensures a specified no. of identical Pods are always running



- Deployment & Statefulness

- ↳ A deployment manages a set of Pods to run an application workload

- ↳ Use Case of Deployment

- ↳ It helps creating new replicas of existing pods

- ↳ If we change a pod then it helps to scale up the new one & gradually scale down the older

- ↳ If a pod crashes, it rolls back to stable state

- ↳ Helps in auto scaling (balances load automatically)

- ↳ ~~Cross~~ Clears up the older state

- Deployment is an abstraction over Pods.

- We work with Deployment

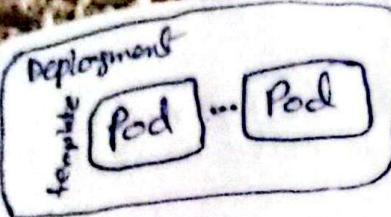
- We can't replicate DBs using Deployment because it would cause data inconsistencies. (all DB pods sharing same Database)

- ↳ This can be done thru StatefulSets

- ↳ It makes sure that DB read/writes are synchronized

- ↳ But it is not easy to work with Statefulsets.

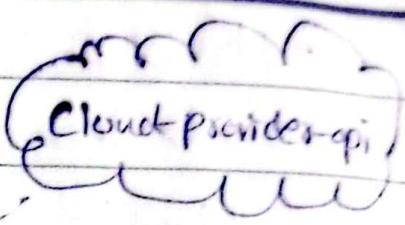
- ↳ recommended to use Statefull apps outside K8s cluster



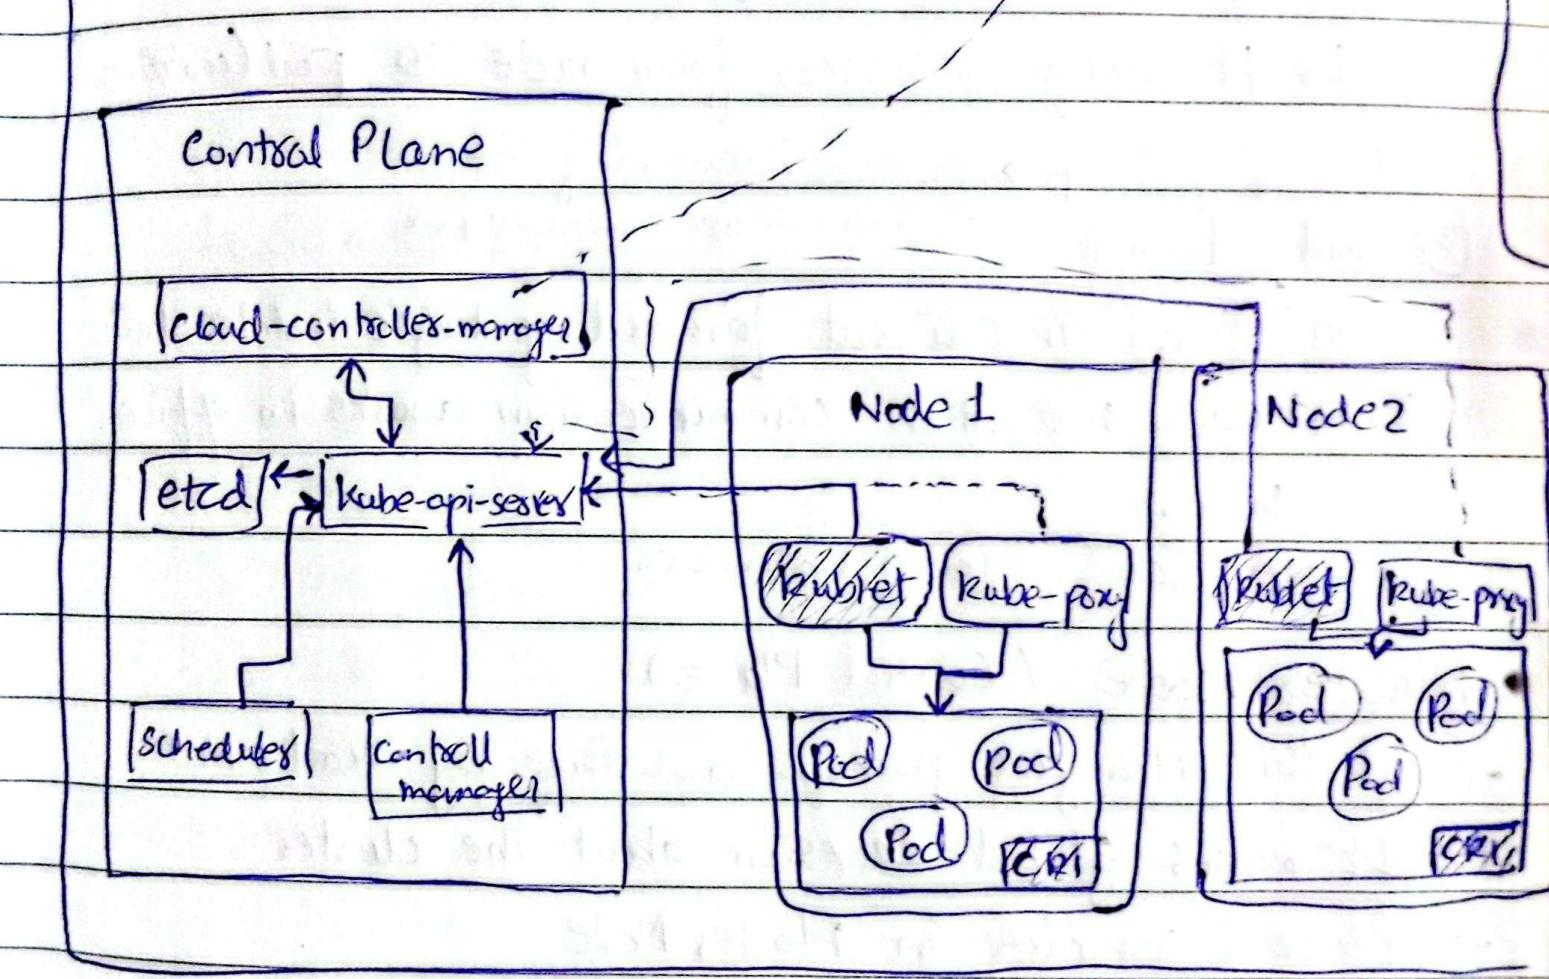
Date _____

Kubernetes Architecture

cluster Nodes



CLOUD CLUSTER



kubernetes Node

- ↳ One of the main component
- ↳ Each node have multiple Pods running on it.
- ↳ Each node have 3 process

(1) Container Runtime → A fundamental component that empowers Kubernetes to run containers
 (Docker Engine) → A software that runs container
 → containerd (default CRT)

Page No. _____

An agent that runs on each node

It takes a set of PodSpecs that are provided through various mechanism and ensures that the containers describe in that Spec is running & healthy

Date _____

② Kubelet

↳ This is the process that schedules the pods

↳ It is the interface to both CRI and the machine (node)

↳ It assigns resources from node to pod (containers)

↳ manages networking

↳ maintains network rules on pod

③ Kube Proxy

↳ It has intelligent forwarding logic inside that makes sure that communication works in efficient way.

↳ Act as a load balancer

• Master Node (Control Plane)

→ All the managing processes are done by Master Node

→ It makes global decisions about the cluster

→ Process running in Master Node

① API-Server → Exposes Kubernetes API

↳ It is like a cluster gateway

↳ It is the frontend for the Kubernetes control plane

↳ The client interacts with API-server.

↳ It gets the initial update

↳ Act as a gatekeeper for authentication

↳ It is designed to scale horizontally

→ a component that watches for newly created pods with no assigned node

Date _____

② Scheduler

- ↳ If we need to add a new pod API-server sends msg to Scheduler.
- ↳ It intelligently identifies on which node to create the Pod
- ↳ Schedules on least busy Node.
- ↳ Schedules only decide on which node to create
- ↳ the actual creation is done on Node by Kubelet
→ Component that runs controller processes

③ Controller Manager

- ↳ If a pod die on a Node so to manage & reschedule it, it is done by this process
- ↳ Detects cluster state changes
- ↳ It request the scheduler to create new pods

④ Etcd

- ↳ Key Value Store of cluster
- ↳ Used as Kubernetes' backing store for all cluster data
- ↳ Etcd is the cluster brain
- ↳ Application data is not stored (only cluster state info is stored)
- ↳ Stores data in a distributed manner

- Master Node is replicated and API-server is load balanced.
- Master Nodes req. less resources than Worker Nodes

→ Types Of Controllers

① Deployment Controller:

- ↳ Ensures the desired no-of Pods are running via Replicaset and manages rolling updates & rollback.
- ↳ Only specifies the no.of pods in a Replicaset.

② ReplicaSet Controller

- ↳ It creates or deletes pods based on the actual and desired no.of pods in Replicaset.

③ Node Controller

- ↳ Monitors node health and marks node as NotReady or removes Pods when nodes fail

④ Job Controller

- ↳ Ensures that a specified no.of Pods successfully complete a one-time task

③ Service Controller

- ↳ Creates and manages the cloud LB for Services of type LoadBalancer

④ Endpoint Controller

- ↳ Update Endpoints or EndpointSlices to map Services to the correct Pod IPs

⑤ Service Account Controller

- ↳ Automatically creates a default ServiceAccount for every namespace and ensures Pods can authenticate with API Server.

"Controller continuously watch the API server and take corrective actions to ensure the actual cluster state always matches the desired state in etcd."

KUBERNETES

Date _____

• Workload Objects

- ↳ Workload objects tell Kubernetes "how" and "where" to run your application containers.
 - ↳ How many copies to run?
 - ↳ On which Nodes?
 - ↳ Whether it runs once or continuously?

→ Types of Workload Objects

① Deployment

- ↳ A deployment is used to run applications that should always be running and can be replaced easily

- ↳ Like a web app manager
- ↳ Runs multiple copies of your app
- ↳ Replaces failed pods
- ↳ Supports rolling updates & rollback
- ↳ Used in: Web servers, APIs, Microservices

↳ In real life we don't make pods we make workload objects and specify all containers inside that then K8s does rest

- Difference b/w Stateful set & deployment.
- In deployment pods have random hashes whereas in stateful set each pod has stable unique ordinal index
 - Pods deleted/created in random order. Pods Date deleted and scaled in predictable order

② Stateful Set

- ↳ For apps that need fixed identity and storage
- ↳ e.g. like a database manager
- ↳ Features
 - ↳ Fixed pod names
 - ↳ stable network identity
 - ↳ Dedicated storage per pod
 - ↳ Identity of pod is preserved when replaced by another

③ DaemonSet

- ↳ A DaemonSet ensures one copy of a Pod runs on every node
- ↳ Automatically runs on a new node
- ↳ No need to define replica count
- ↳ Used for log collector or monitoring agent

④ Job

- ↳ A Job runs a task until it finishes successfully
- ↳ What it does
 - ↳ Runs to completion
 - ↳ Retries on failure
 - ↳ Stops when done

⑤ CronJob

- ↳ A CronJob runs Jobs on a schedule
- ↳ Used cases: Nightly backups, Cleanup jobs, Report Generation

• Namespace

- ↳ A logical isolation mechanism in Kubernetes that allows multiple teams or env to share the same cluster while keeping resources organized, secure and manageable.
- ↳ In traditional setups we have diff servers for diff env like dev, test, pat, prod etc. So in order to achieve this in kubernetes we use "namespace".
- ↳ Each name space runs its own version of the application, database configurations and ingress routing.

• Volume

- ↳ An abstraction that let us persist data
- ↳ Necessary bcz pods are ephemeral - meaning data is deleted when pod is deleted

• Benefits Of Kubernetes

- ① Scale the application as needed from commandline or UI
- ② Automated rollouts and rollbacks
- ③ Containers get their own IP so you can put a set of containers behind a single DNS name for loadbalancing
- ④ Automatically mount local or public cloud or a network storage
- ⑤ Create and update secrets & configs without image rebuilding
- ⑥ Self Healing
 - ↳ Restarting failed containers
 - ↳ Replacing & Rescheduling containers as nodes die
 - ↳ Killing containers that don't respond
- ⑦ Automatic resource scheduling based on requirement & constraints.

• Kubectl

↳ Command line tool used to:

- Talk to the Kubernetes API Server
- Create, update, delete and inspect Kubernetes objects

① Check cluster connectivity

» kubectl cluster-info

② Get Commands

» kubectl get <resource>

↳ resource can be nodes, pods, services, deployments
namespaces

↳ get detailed output

» kubectl get pods -o wide

③ Describe commands .

↳ Shows detail information

» kubectl describe <resource> <resource_name>

pod

my-pod

Page No. □

Date

↳ Used when pod is crashing, image is not pulling,
health checks fail

④ Create and Apply Commands

↳ Create from Yaml

⇒ kubectl apply -f my-deployment.yaml

↳ Creating Imperatively

⇒ kubectl create deployment nginx --image=nginx

⑤ Delete Commands

⇒ kubectl delete pod my-pod

⇒ kubectl delete -f my-deployment.yaml

⇒ kubectl delete all --all -n dev

⇒ delete everything in namespace

⑥ Logs & Exec (Debugging)

⇒ kubectl logs my-pod

→ log from specific container

⇒ kubectl logs my-pod -c my-container

Date _____

↳ Get shell inside container

» kubectl exec -it my-pod -- /bin/bash

⑦ Namespace

» kubectl get ns → list all namespaces

↳ Create namespace

» kubectl get create ns dev

⑧ Scaling

» kubectl scale deployment myapp --replicas=5

⑨ Edit Deployment

» kubectl edit deployment my-deployment

Date

• KUBERNETES CONFIGURATION FILE (YAML)

→ Deployment File

apiVersion: apps/v1

kind: Deployment

→ defines the object we are creating
like Service, StatefulSet, Daemon...

metadata:

name: myapp-deployment

} used for identification, label

labels:

} Selection. "name" must be unique
in a namespace

app: myapp

spec:

→ write the specification of deployment

replicas: 3

→ no. of pods to run

selector:

} deployment manages pods with this
label

matchLabels:

} Must match pod template label
exactly

app: myapp

Date _____

template :

metadata :

labels :

app : myapp

spec : → spec for
a pod

containers:

- name : myapp-container

image : nginx:1.25

must match the selector label

container is an array therefore
we can multiple containers

ports :

- containerPort : 8080

env.:

- name : ENV

value : "dev"

Date

→ Service File

apiVersion: v1

kind: Service

metadata:

name: myapp-service

spec: → Can be NodePort ⇒ Expose via Node IP

Type: ClusterIP

LoadBalancer ⇒ Cloud LB (AWS, GCP)

ExternalName ⇒ DNS mapping

ClusterIP ⇒ Internal only (default)

Selector:

app: myapp → ~~also~~ match the pod
label

should

ports:

- port: 80

→ port on which
client request

client

↓ 80

targetPort: 8080

myapp-service

↓ 8080

→ port on which service
will request to pod

Pod

should match "containerPort"
in Pod specification.

optional, if not define default value = port

Page No. □

Date _____

→ ConfigMap File

apiVersion: v1

kind: ConfigMap

metadata:

name: my-configmap

data:

BACKEND_URL: https://myapp.com/api/1:5000

DB_HOST: my-sql-service

DB_PORT: "3306"

→ Secret File

apiVersion: v1

kind: ConfigMap

metadata:

name: my-appsecret

Type: Opaque

data:

↗ base64 encoded

DB_USER: abged2zt= #admin

DB_PASSWORD: googlm34= #pass123

Date

- Usage:

↳ Used in pod's specification

spec:

containers:

- name: mysql

image: mysql

env:

- name: MYSQL_USER

valueFrom:

secretKeyRef:

name: DB_PASSWORD my-app-secret

key: DB_PASSWORD

- name: MYSQL_PASSWORD

:

Ref

ConfigMap use kya hoga to bas "secretKeyRef" →
"configMapKeyRef".

Date _____

→ StatefulSet File

↳ Everything is similar to deployment file only we add "volume" here.

```
apiVersion: apps/v1
```

```
kind: StatefulSet
```

```
metadata:
```

```
  name: mysql
```

```
spec:
```

```
  serviceName: my-sql-service → required
```

```
  replicas: 1
```

```
  selector:
```

```
    matchLabels:
```

```
      app: mysql
```

```
template:
```

```
  metadata:
```

```
    labels:
```

```
      app: mysql
```

Date

spec:

containers:

```
- name: mysql-container  
  image: mysql
```

env:

```
- name: MYSQL_ROOT_PSWD
```

valueFrom:

secretKeyRef:

```
      name: mysql-secret
```

```
      key: db-password
```

```
- name: MYSQL-USER
```

ports:

```
- containerPort: 3306
```

VolumeMounts:

```
- name: mysql-data  
  mountPath: /var/lib/mysql
```

Date _____

!!

VolumeClaimTemplates:

- metadata:

name: mysql-data

Spec:

accessMode: ["ReadWriteOnce"]

resources:

requests:

storage: 10Gi

→ Headless Service

↳ Stateful set requires headless service

apiVersion: v1

kind: StatefulSet

service:

name: mysql-service

Spec:

clusterIP: None

Headless

selector:

app: mysql

ports:

- port: 3306