

# KUBERNETES

Date \_\_\_\_\_

- Open source, container orchestration tool
- Helps you manage containerized applications in different environment

The need for a container orchestration tool

- Trend from Monolithic to microservice architecture
- Increase usage of containers
- Demand for a proper way of managing those hundreds of containers

## Features

- ↳ High availability / No downtime
  - ↳ Scalability or high performance
  - ↳ Disaster recovery - backup and restore
- Through Kubernetes we can run our app on any cloud provider. It provides an abstraction layer for us.


Page No.

Date \_\_\_\_\_

## ► Main Components of Kubernetes

### • Pod

- ↳ Smallest unit of K8s
- ↳ Abstraction over containers
- ↳ It contains one or more application containers that are tightly coupled
- ↳ But mostly we should run one-container per Pod
- ↳ It creates a layer or running env on top of containers.
- ↳ K8s has made this abstraction so that we are not dependent on certain specific containerization technology.
- ↳ Each Pod gets its own IP Address (not the container!)
- ↳ Pods can communicate with each other using IP
- ↳ Pods can die easily, so when they created again they get new IP. So if two pods communicate each other then we should update the IP.



### • Service

- ↳ For this reason service is used
- ↳ It is basically a permanent IP Address that can be attached to each Pod.

Page No. \_\_\_\_\_

→ Service & also a load balancer for Pod. Like we have two Pods of AddToCart Service. So the service sends request to the Pod.

Date \_\_\_\_\_

↳ Lifecycle of Pod & service not connected.

↳ Even if Pod dies it will not loose its permanent IP Addr.


→ External Service (like we need to run an app on browser (easily accessible))

→ Internal Service (like DBs)

## Ingress

→ Exposes HTTP & HTTPS routes from outside the cluster to services within the cluster.

→ Traffic routing is controlled by rules defined on the ingress resource.



- Config Map

↳ Used to store non-confidential data in key value pairs.

↳ Pods can use them as env variables.

↳ ConfigMap allows you to decouple env-specific configurations from your container images.

↳ ConfigMaps can also be mounted as data volumes.

- Secret

↳ It contains small amount of sensitive data such as a password / a token or a key.

↳ They are similar to ConfigMap.

↳ They store data in base64 encoding format.

- Volumes

↳ By default there is no data persistence on restarting of Pod.

↳ But can be done ~~not~~ thru Volume.

↳ This volume is separate from cluster.

↳ It can be the part of the local machine or remote.

**Replicaset:** Is a Kubernetes controller that ensures a specified no. of identical Pods are always running

→ Pod Blueprint

Date \_\_\_\_\_

- Deployment & Statefulness

- ↳ A deployment manages a set of Pods to run an application workload

- ↳ Use case of Deployment

- ↳ It helps creating new replicas of existing pods

- ↳ If we change a pod then it helps to scale up the new one & gradually scale down the older

- ↳ If a pod crashes, it rolls back to stable state

- ↳ Helps in auto scaling (balances load automatically)

- ↳ Deletes Clears up the older state

- Deployment is an abstraction over Pods.

- We work with Deployment

- We can't replicate DBs using Deployment

- because it would cause data inconsistencies (all DB pods sharing same Database)

- ↳ This can be done thru StatefullSets

- ↳ It makes sure that DB read/writes are synchronized

- ↳ But it is not easy to work with Statefullsets.


- ↳ recommended to use Statefull apps outside K8s cluster

Page No. \_\_\_\_\_

Date

## • Kubernetes Architecture

### Kubernetes Nodes



### • kworker Node

- ↳ One of the main component
- ↳ Each node have multiple Pods running on it.
- ↳ Each node have 3 process

- ① Container Runtime  
(Docker Engine)

Page No.

## ② Kubelet

↳ This is the process that schedules the pods

↳ It is the interface to both CRI and the machine (node)

↳ It assigns resources from node to pod (containers)

## ③ Kube Proxy

↳ It has intelligent forwarding logic inside that makes sure that communication works in efficient way.

## • Master Node (Control Plane)

→ All the managing processes are done by Master Node

→ It makes global decisions about the cluster

→ Processes running in Master Node

## ① API-SERVER

↳ It is like a cluster gateway

↳ It is the frontend for the Kubernetes control plane

↳ The client interacts with API-server.

↳ It gets the initial update

↳ Act as a gatekeeper for authentication

## ② Schedules

- ↳ If we need to add a new pod Api-server sends req to Schedules.
- ↳ It intelligently identifies on which node to create the Pod
- ↳ Schedules on least busy Node.
- ↳ Schedules only decide on which node to create
- ↳ the actual creation is done on Node by Kubelet

## ③ Controller Manager

- ↳ If a pod die on a Node so to manage & reschedule it, it is done by this process
- ↳ Detects cluster state changes
- ↳ It request the schedules to create new pods

## ④ Etcd

- ↳ Key value store of cluster
- ↳ Used as Kubernetes backing store for all cluster data
- ↳ Etcd is the cluster brain
- ↳ Application data is not stored (only cluster state info is stored)

Date \_\_\_\_\_

- Master Node is replicated and API-Server is load balanced.
- Master Nodes req. less resources than Worker Nodes